

POPEN (Pipe Open)

1. If you need to run a program or any general shell command from C code, and either read its output or send data to its standard input, you can use **popen**.
2. Write a program that generates the "99 bottles of beer" song, and displays it in a pager like less.

```
int main(int argc, char** argv) {
    int i;
    FILE* fp;

    fp = popen("less", "w");
    for(i=99; i>0; i--) {
        fprintf(fp, "%d bottles of beer on the wall,\n", i);
        fprintf(fp, "    %d bottles of beer.\n", i);
        fprintf(fp, "If one of those bottles should happen to fall,\n");
        fprintf(fp, "    %d bottles of beer on the wall.\n\n", i-1);
    }
    pclose(fp);
    return 0;
}
```

3. Write a C program that displays the user with the most processes in the system, and how many processes he or she has. Use a Shell command to find the data.

```
int main(int argc, char** argv) {
    int n;
    char u[64];
    FILE* fp;

    fp = popen("ps -ef | awk '{print $1}' | sort | uniq -c | sort -nr | head -n 1", "r");
    fscanf(fp, "%d %s", &n, u);
    pclose(fp);
    printf("%s: %d\n", u, n);
    return 0;
}
```

UNIX PROCESS FILE DESCRIPTOR MANIPULATION WITH DUP() AND DUP2()

1. The goal of this section is to see behind the scenes of input/output redirections, and eventually be able to write in C the equivalent of the command `ps -ef | grep -E "^root" | awk '{print $2}'`
2. Every process in a Unix-like operating system has a **file descriptor table** that is just a small table (array) that keeps track of all the **open files** (and other things like pipes or sockets) that the process is using.

Example :

```
int main(int argc, char**argv) {
    int fd, myfifo, pa[2], pb[2];

    fd = open("a.txt", O_RDWR);
    myfifo = open("myfifo", O_RDONLY);
    pipe(pa);
    pipe(pb);

    return 0;
}
```

The file descriptor table will look as follows

Index	Value
0	Handle for reading from the console (standard input)
1	Handle for writing to the console (standard output)
2	Handle for writing to the console (standard error)
3	Handle for reading/writing to file a.txt. Variable <code>fd</code> has the value of the index, i.e. 3
4	Handle for reading from FIFO myfifo. Variable <code>myfifo</code> has the value of the index, i.e. 4
5	Handle for reading from the first PIPE created. Variable <code>pa[0]</code> has the value of the index, i.e. 5
6	Handle for writing to the first PIPE created. Variable <code>pa[1]</code> has the value of the index, i.e. 6
7	Handle for reading from the second PIPE created. Variable <code>pb[0]</code> has the value of the index, i.e. 7
8	Handle for writing to the second PIPE created. Variable <code>pb[1]</code> has the value of the index, i.e. 8

4. The file descriptor table can be manipulated using the system calls **dup()** and **dup2()**.
- int dup(int oldfd)** makes a copy of the handle at index **oldfd** to a new entry, and returns the index of the entry just created. For example, adding line **int x = dup(myfifo)** to the program above, just before the **return**, will result in the following line being added to the file descriptor table, and variable **x** will have the value 9.

9	Handle for reading from FIFO myfifo.
---	--------------------------------------

- int dup2(int oldfd, int newfd)** makes a copy of the handle at index **oldfd** to index **newfd**, overwriting the existing value (it also closes it silently before overwriting it). For example, adding line **dup2(p[0], 0)** to the program above, just before the **return**, will result in line 0 containing the handle for reading from the first PIPE created. Consequently, any reading from the standard input will use the first PIPE created because by using **dup2(p[0], 0)**, you're telling the **grep** process: 'Don't read from the keyboard (stdin = 0) anymore; instead, read from the read end of PIPE'.

0	Handle for reading from the first PIPE created.
---	---

5. Let's implement in C a program that does the equivalent of running the command **ps -ef | grep -E "^root" | awk '{print \$2}'**, by running the three programs and transferring the data between them using pipes.

```
int main(int argc, char** argv) {
    int p2g[2], g2a[2];

    pipe(p2g); pipe(g2a);
    if(fork() == 0) {
        close(p2g[0]); close(g2a[0]); close(g2a[1]);
        dup2(p2g[1], 1);
        execlp("ps", "ps", "-ef", NULL);
        exit(1);
    }
    if(fork() == 0) {
        close(p2g[1]); close(g2a[0]);
        dup2(p2g[0], 0); dup2(g2a[1], 1);
        execlp("grep", "grep", "-E", "^root", NULL);
        exit(1);
    }
}
```

```
if(fork() == 0) {
    close(p2g[0]); close(p2g[1]); close(g2a[1]);
    dup2(g2a[0], 0);
    execlp("awk", "awk", "{print $2}", NULL);
    exit(1);
}

close(p2g[0]); close(p2g[1]);
close(g2a[0]); close(g2a[1]);

wait(0); wait(0); wait(0);

return 0;
}
```

6. How do you undo a **dup2** call? Use **dup** to make a copy of the entry that will be overwritten, and when you want to reset it, use again **dup2** with the copy.

```
int x = dup(1);
dup2(p[1], 1);
....
dup2(x, 1);
```

UNIX IPC SHARED MEMORY

- UNIX IPC is a set of mechanisms (other than pipe and FIFO) for communicating among processes (IPC = Inter-Process Communication)
 - Semaphore** : synchronization mechanism
 - Message queues** : message-based communication
 - Shared memory** : communication through a common memory region

Just like FIFOs, IPCs allow any two processes to communicate through them, but unlike FIFOs, IPCs are not files on disk. They are identified by a number that is unique in the system, this number is chosen by the developer and provided to all processes that need to communicate through that specific IPC.

- IPC cleanup
 - IPCs are persistent structures in the operating system
 - The operating system imposes limits for the size and number of IPCs
 - If not cleanup properly, the system will refuse to allow the creation of new IPCs
 - You can see all existing IPCs using the command **ipcs**
 - You can delete an IPC using command **ipcrm**, as long as you have permissions on that IPC

> **ipcs**

----- Message Queues -----

key	msqid	owner	perms	used-bytes	messages
-----	-------	-------	-------	------------	----------

----- Shared Memory Segments -----

key	shmid	owner	perms	bytes	nattch	status
0x58df7f12	0	rares	644	100	0	

----- Semaphore Arrays -----

key	semid	owner	perms	nsems
-----	-------	-------	-------	-------

> **ipcrm -M 0x58df7f12**

3. Shared memory API

- Create or get a handle to an existing shared memory segment: **shmget**
- Attach - map the shared memory segment to a pointer in your process: **shmat**
- Detach - unmap the shared memory segment from you pointer: **shmdt**
- Control (configure, delete, etc.) the shared memory segment: **shmctl**